


Alexnet CNN

CS231n

Stanford

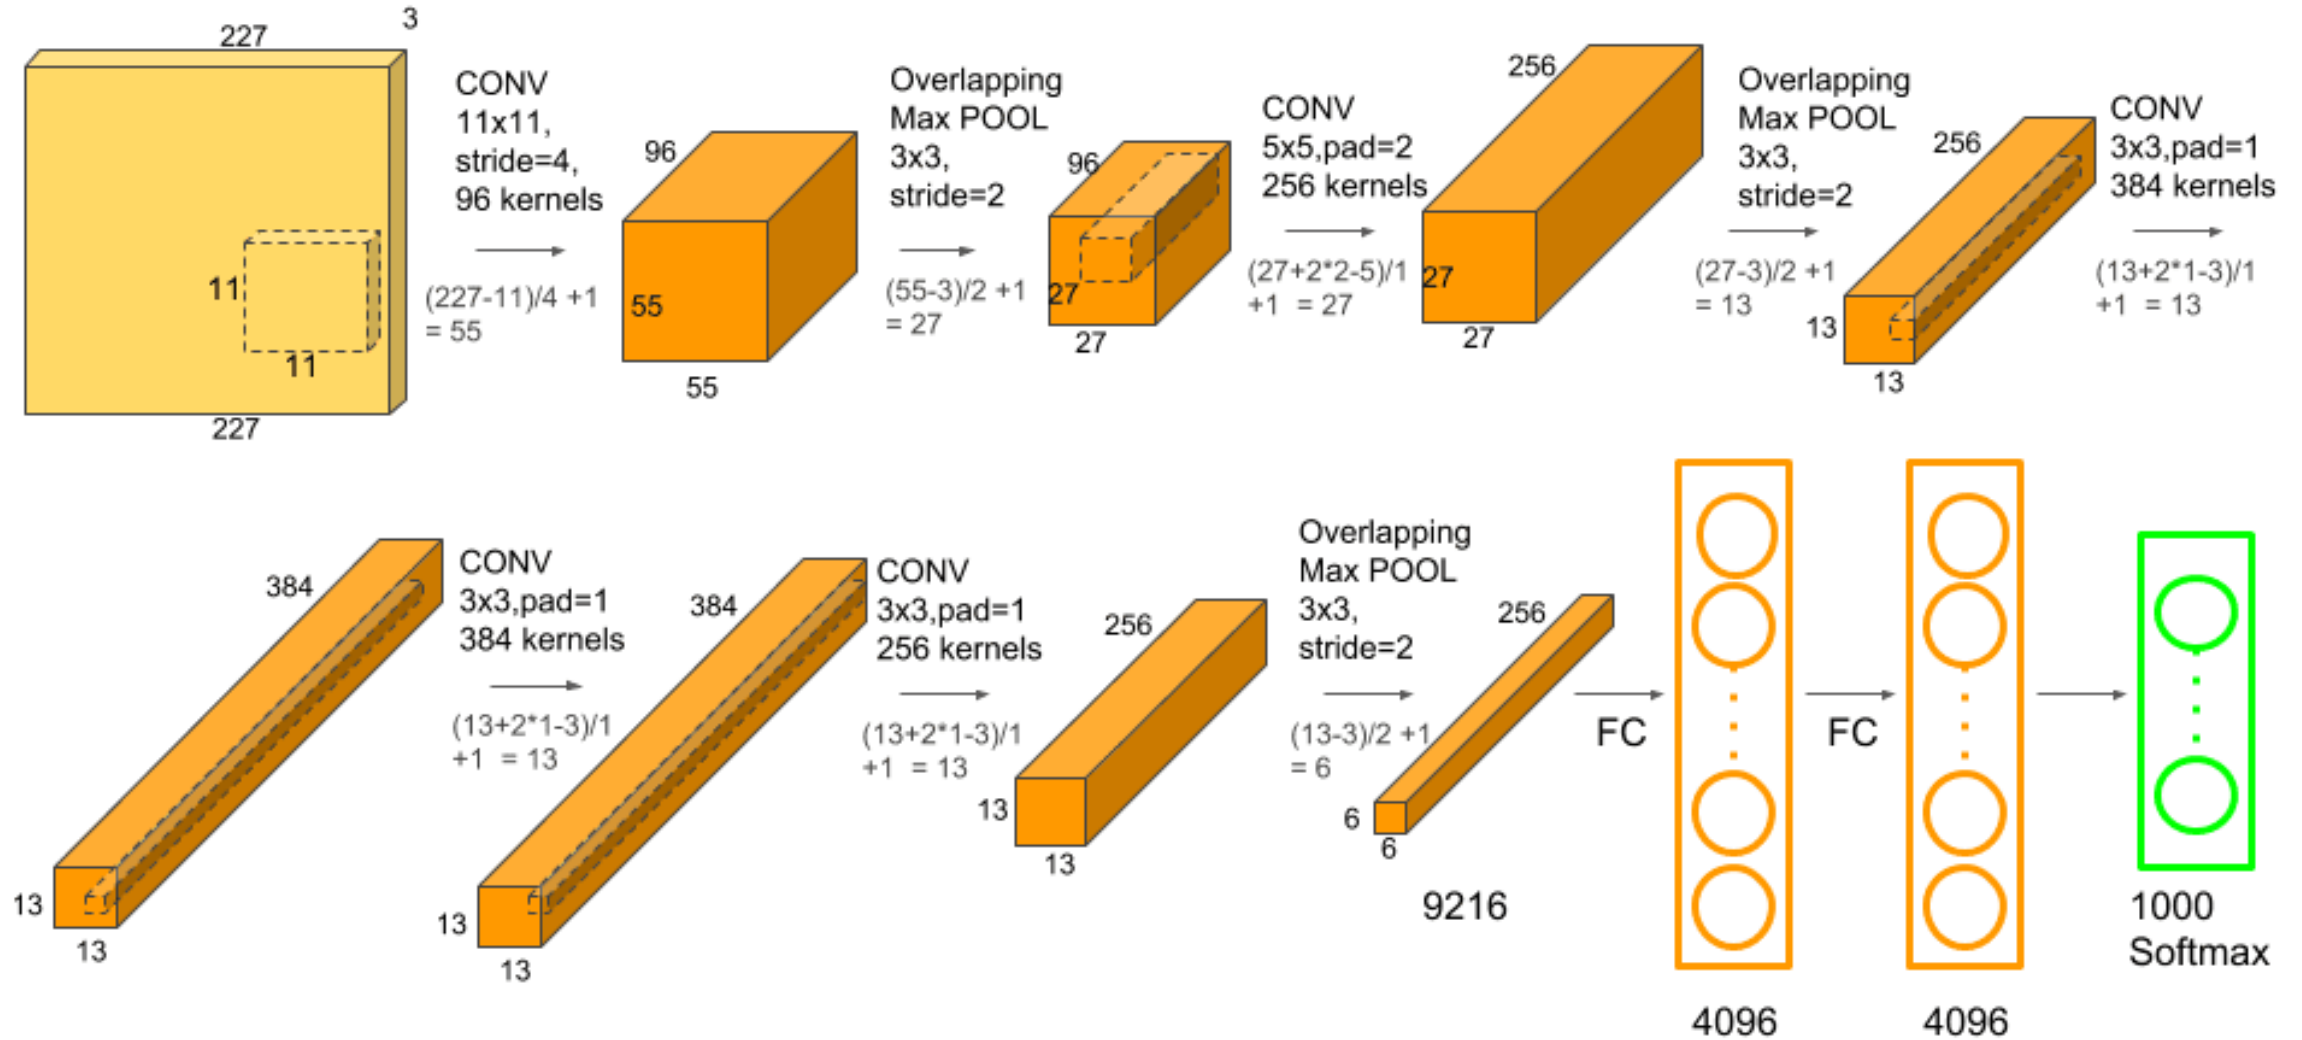


Paper Outcomes

1. The ReLU non-linearity is applied to the output of every convolutional and fully-connected layer.
2. We applied this normalization after applying the ReLU nonlinearity in certain layers.
3. If we set $s < z$, we obtain overlapping pooling. This is what we use throughout our network, with $s = 2$ and $z = 3$.
4. The first convolutional layer filters the $224 \times 224 \times 3$ input image with 96 kernels of size $11 \times 11 \times 3$ with a stride of 4 pixels
5. The second convolutional layer takes as input the (response-normalized and pooled) output of the first convolutional layer and filters it with 256 kernels of size $5 \times 5 \times 48$.
6. The third, fourth, and fifth convolutional layers are connected to one another without any intervening pooling or normalization layers. The third convolutional layer has 384 kernels of size $3 \times 3 \times 256$ connected to the (normalized, pooled) outputs of the second convolutional layer.
7. The fourth convolutional layer has 384 kernels of size $3 \times 3 \times 192$, and the fifth convolutional layer has 256 kernels of size $3 \times 3 \times 192$. The fully-connected layers have 4096 neurons each.
8. We use dropout in the first two fully-connected layers [...]

Important Formula
$$= \frac{(n+2p-f)}{s} + 1,$$

 n = size of image, p = padding, f = filter, s = stride



Network architecture

- The network consists of 5 *Convolutional* layers and 3 *Fully Connected* Layers
- *Max Pooling* is applied Between the layers:
 - 1conv-2conv
 - 2conv-3conv
 - 5conv-1fc
- Before Max Pooling a normalization technique is applied.